

TEJAS MISTRY

(315) 849-8584 | tejas2000930@gmail.com | linkedin.com/in/tamistry | tamistry.netlify.app |

SUMMARY

Data Engineer with 2+ years of experience across hospitality, entertainment, and academic domains, building and optimizing batch ETL/ELT pipelines, cloud data warehouses, and distributed processing systems on AWS. Reduced pipeline refresh latency by 75% through Snowpipe and Airflow-orchestrated Glue workflows. Hands-on with Snowflake, dbt, PySpark, AWS Glue, Redshift, and Airflow delivering reliable, analytics-ready data pipelines for BI reporting.

SKILLS

Languages: Python, SQL, R
Data Engineering: ETL/ELT Pipelines, Data Modeling, dbt, Apache Airflow, Data Quality Validation, Batch Processing
Processing & Warehousing: PySpark, Apache Spark, Databricks, Snowflake, Amazon Redshift, AWS Athena, Apache Kafka
Cloud Platforms: AWS (Glue, Lambda, S3, Snowpipe), Azure
Visualization & BI: Power BI (DAX, Power Query), Tableau, Looker Studio, DOMO
Tools: Docker, Git, CI/CD, PostgreSQL, Parquet, JupyterLab

EXPERIENCE

Data Engineer, Statistical Research | *Syracuse University* | Syracuse, NY Jul 2025 - Present

- Build Python/SQL ELT pipelines orchestrated via Apache Airflow, using AWS Glue and Lambda to ingest and transform 120K+ academic records, survey responses, and enrollment data into Amazon Redshift.
- Implement SQL-based schema validation across 6 Redshift tables, enforcing null checks and type consistency to catch data quality failures before downstream analytics delivery.
- Query and stage intermediate datasets in AWS Athena against S3-stored files to serve ad hoc research requests without triggering full Redshift compute loads.
- Configure and maintain DOMO integration with Redshift to automate reporting on student enrollment trends and performance metrics, eliminating manual data pull requests from research teams.
- Containerize Airflow pipeline workflows using Docker and enforce Git-based pull request reviews with CI/CD practices to maintain deployment-ready pipeline code.

Data Science Intern | *Bandhouse Music Group Co* | Nashville, TN May 2024 - Sep 2024

- Engineered Python-based data transformation pipelines to aggregate data from 5+ streaming and social platforms into unified analytical datasets, improving visibility across artist performance KPIs.
- Designed reusable structured datasets standardizing artist engagement metrics (playlist-add velocity, stream-to-follower growth, engagement ratio) for consistent campaign analytics delivery.
- Built regression and clustering models to forecast listener growth and segment audiences for promotional planning.
- Ran A/B tests and statistical analyses (t-tests, chi-square) on 3,000+ user interactions, delivering a statistically significant 6% incremental lift in user retention at 95% confidence.

Jr. Data Engineer | *J&J Hospitality Pvt. Ltd.* | Mumbai, India May 2022 - Jul 2023

- Designed and maintained batch ETL pipelines orchestrated via Apache Airflow DAGs, coordinating AWS Glue and Lambda to ingest purchasing, supplier, sales, and inventory data into Snowflake, reducing data refresh latency from 4 hours to 1 hour.
- Built and optimized data ingestion workflows using Amazon S3 and Snowpipe across 4 source domains, reducing ingestion latency from 90 minutes to 10 minutes and accelerating data availability for 3 downstream BI analysts.
- Implemented star-schema data models using dbt, defining fact and dimension tables to support scalable analytics and deliver curated datasets powering operational KPI dashboards.
- Developed PySpark-based transformations within AWS Glue to clean, standardize, and validate 680,000+ structured records, ensuring schema consistency and high data quality for downstream analytics.
- Wrote complex SQL queries with CTEs and window functions in Snowflake to accelerate data validation workflows and improve analytical query performance.
- Delivered Power BI dashboards using DAX and Power Query on Snowflake data, enabling real-time monitoring of inventory, supplier costs, and operational KPIs.

PROJECTS

Airfare Prediction and Optimization PySpark

Python | AWS S3 | MLlib | Pandas | Gradient Boosting | Random Forest | Jupyter

- Designed PySpark ETL pipelines on AWS S3 processing 2M+ flight records, converting raw data to Parquet for optimized storage and query performance, with feature engineering workflows enabling quality inputs for ML-based airfare forecasting.
- Implemented Gradient Boosting and Random Forest models achieving low prediction error (RMSE: 12.6), delivering a data-driven pricing tool for airfare revenue optimization.

vBay Auction Data Warehouse & Business Intelligence

AWS S3 | AWS Glue | Snowflake | dbt | SQL | Power BI | PostgreSQL

- Architected cloud-native Medallion Architecture (Bronze/Silver/Gold) on AWS S3, Glue, and Snowflake with dbt star-schema models to centralize auction platform data (bids, transactions, users) for scalable querying.
- Built interactive Power BI dashboards surfacing real-time insights into bidding trends, category performance, and regional engagement to support optimized pricing strategies.

EDUCATION

Syracuse University | M.S. in Applied Data Science | Syracuse, NY Aug 2023 - May 2025

University of Mumbai | B.E. in Computer Engineering | Mumbai, India Aug 2018 - May 2022