

Tejas Mistry

(315) 849-8584 | tejas2000930@gmail.com | [linkedin.com/in/tamistry](https://www.linkedin.com/in/tamistry) | tamistry.netlify.app

SUMMARY

Data Engineer with 2+ years building production-grade ETL/ELT pipelines and dimensional data models across hospitality, entertainment, and academic domains. Reduced pipeline refresh latency by 75% at J&J Hospitality and delivered analytics-ready datasets across 2M+ records at Syracuse University. Hands-on with Snowflake, dbt, Airflow, Microsoft Fabric, and AWS, currently extending into LLM-integrated data systems via LangGraph and Claude API.

EXPERIENCE

Data Engineer, Statistical Research | *Syracuse University* | Syracuse, NY Jul. 2025 – Present

- Maintain and enhance 6 Microsoft Fabric Data Factory ELT pipelines integrating PeopleSoft SIS, Blackboard LMS, and administrative systems into OneLake, processing 2M+ historical records into a shared analytics foundation.
- Develop SQL and Python transformations in Fabric notebooks to standardize enrollment, course, and term data, resolving duplicates, missing values, and inconsistent field types before publishing validated Silver-layer datasets.
- Design 10+ Gold-layer fact and dimension tables, standardizing enrollment and course metrics for Power BI semantic models and letting stakeholders reuse governed logic instead of rebuilding transformations per report.
- Improve production pipeline reliability by diagnosing upstream schema and data-type changes, correcting transformation logic, and implementing validation checks to catch incompatible records before they reach downstream institutional reports.

Data Science Intern | *Bandhouse Music Group Inc* | Nashville, TN May 2024 – Sep. 2024

- Centralized KPI data across 5 platforms (Spotify, Instagram, TikTok, YouTube, Apple Music) for 40+ artists into unified Python transformation pipelines, enabling weekly cross-platform reporting across 2M+ monthly data points.
- Built regression and clustering models to forecast listener growth and segment audiences for a 180K MAU platform, informing artist marketing and release-timing decisions.
- Drove a 6% incremental lift in user retention (95% confidence) via A/B testing across 3,000+ listener interactions.

Jr. Data Engineer | *J&J Hospitality Pvt. Ltd.* | Mumbai, India May 2022 – Jul. 2023

- Designed and maintained Airflow-orchestrated batch ETL pipelines coordinating AWS Glue and Lambda to ingest purchasing, supplier, sales, and inventory data into Snowflake, reducing pipeline refresh latency by 75% (4 hrs to 1 hr).
- Built near-real-time data ingestion pipelines using Amazon S3 event triggers and Snowpipe, reducing ingestion latency from 90 minutes to 10 minutes across 4 source domains for downstream BI consumers.
- Implemented dbt star-schema models (8 fact and dimension tables) over Snowflake, enabling analysts to build KPI dashboards without engineering support.
- Processed and validated 680,000+ structured records using PySpark transformations within AWS Glue, enforcing schema consistency and data governance standards.
- Wrote complex SQL with CTEs and window functions in Snowflake for data validation and query performance tuning.
- Built Power BI dashboards (DAX, Power Query) on Snowflake with a 15-minute scheduled refresh cycle, enabling near-real-time monitoring of inventory, supplier costs, and operational KPIs for 3 BI analysts and stakeholders.

PROJECTS

LLM-Powered Analytics Agent | *LangGraph, FastAPI, Claude API, PostgreSQL, Docker*

- Built an NL-to-SQL agent using LangGraph and the Claude API against a multi-table PostgreSQL warehouse, reducing ad hoc SQL turnaround from hours to seconds and achieving 91% query accuracy on a 200-query benchmark.
- Implemented human-in-the-loop checkpoints for query validation, preventing erroneous writes and ensuring data governance.

Airfare Prediction and Optimization | *Databricks, PySpark, Spark MLlib, Statistical Modeling, Feature Engineering*

- Modeled airfares from 2M+ flight records using PySpark on Databricks, benchmarking GBRegressor, RandomForest, Linear, and Decision Tree models, with the top model reaching 0.99 R-squared and predicting fares within a \$13 margin.
- Addressed data quality and dimensionality issues by engineering 15+ features and applying KNN imputation, IQR outlier detection, and correlation analysis, validating results via k-fold cross-validation to quantify seasonality.

SKILLS

Languages: Python, SQL, R

Data Engineering: ETL/ELT Pipelines, Batch & Streaming, Data Modeling, Data Warehousing, dbt, Apache Airflow, Apache Kafka, Data Governance, SLA Management

Processing & Warehousing: Snowflake, PySpark, Apache Spark, Databricks, PostgreSQL, Star-Schema, Medallion Architecture

Cloud & DevOps: Microsoft Fabric, AWS (Glue, Lambda, S3, Redshift), Azure, Docker, CI/CD, GitHub Actions, Terraform

BI & Visualization: Power BI (DAX, Power Query), Tableau, Looker Studio, DOMO

AI & LLM: LangGraph, LangChain, Claude API, RAG, FAISS, Vector Database, NL-to-SQL, FastAPI, Prompt Engineering

EDUCATION

Syracuse University | *M.S. Applied Data Science*

Syracuse, NY | *Aug. 2023 – May 2025*

University of Mumbai | *B.E. Computer Engineering*

Mumbai, India | *Aug. 2018 – May 2022*